



Team: The Night's Watch

ReasonAudit: Evaluating Error Awareness and Self-Correction in Large Language Models

Department of Computer Science and Engineering, United International University

Supervisor

Sadia Islam, Assistant Professor

Co-Supervisor

Mr. Nahid Hossain, Assistant Professor

Course Teacher

Dr. Riasat Azim, Associate Professor

Md. Assaduzzaman Nur (0112230442)


Shahriar Yasin (0112230638)

Shakib Ahmed (0112230421)

Sadid Ahmed (0112330154)

Rukan Mia (0112231036)

Can an LLM catch its own mistakes, or does second-guessing make things worse?



We measure how often self-review fixes errors, and how often it ruins correct answers.


Motivation

- LLMs are used in coding, teaching, and research, often with little human checking.
- They can give wrong answers with full confidence.
- Self-review is fragile: models may miss errors, keep bad answers, or change correct answers to wrong ones.
- We need a precise measure of when self-correction helps and when it hurts.

Objectives

- Build a labeled reasoning dataset with verified answers and difficulty tags.
- Measure Correction Rate, Damage Rate, and accuracy change after self-review.
- Compare open-weight and proprietary models across reasoning types.
- Link difficulty factors to self-correction success.

Methodology



Controls: neutral review · independent second attempt

Models Claude Haiku 4.5 GPT-6 Luna Gemini 3.8 Flash LLaMA 3 8B

Dataset


Mathematics Formal Logic Commonsense

Difficulty modifiers each question tagged yes / no

Ambiguity Missing info. Contradiction Distracting info.

- Domains: Mathematics · Formal Logic · Commonsense
- Difficulty modifiers (yes/no tag per question): Ambiguity · Missing information · Contradiction · Distracting information
- Built by: collecting questions from public datasets, then manual annotation and tagging. Each question keeps its source and a flag if modified.
- Scale: 20-question pilot, designed to scale to 2,000.

Outcome Buckets




Metrics

Correction Rate $\frac{WC}{WC + WW}$ wrong answers fixed

Damage Rate $\frac{CW}{CC + CW}$ correct answers ruined


- Correction Rate = $WC / (WC + WW)$: share of wrong answers the model fixed
- Damage Rate = $CW / (CC + CW)$: share of correct answers the model ruined
- Accuracy Change = Final Accuracy – Initial Accuracy
- Analysis: McNemar's test, 95% confidence intervals, breakdown by domain and difficulty tag.

Research Gap




Pilot System

- Two-stage pipeline in Google Sheets + Apps Script, with a web pilot tool.
- Logs every prompt and raw response for traceability.
- Automatic retry on rate limits, CSV export.



Status and Next Steps



References

- Huang et al., Large Language Models Cannot Self-Correct Reasoning Yet, ICLR 2024
- Tyen et al., LLMs Cannot Find Reasoning Errors, but Can Correct Them Given the Error Location, ACL Findings 2024
- Tsui, Self-Correction Bench, COLM 2026
- Madaan et al., Self-Refine, NeurIPS 2023
- Yuan et al., Hidden Error Awareness in Chain-of-Thought Reasoning, 2026